

1 Linking assumptions and number of response options modulate inferred scalar implicature  
2 rate

3 Masoud Jasbi<sup>1</sup>, Brandon Waldon<sup>1</sup>, & Judith Degen<sup>1</sup>

4 <sup>1</sup> Stanford University

5 Author Note

6 Add complete departmental affiliations for each author here. Each new line herein  
7 must be indented, like this line. Enter author note here.

8 Correspondence concerning this article should be addressed to Masoud Jasbi, Postal  
9 address. E-mail: [my@email.com](mailto:my@email.com)

## Abstract

10

11 Enter abstract here. Each new line herein must be indented, like this line.

12       *Keywords:* scalar implicature; methodology; linking assumption; experimental  
13 pragmatics; truth-value judgment task

14       Word count: X

Linking assumptions and number of response options modulate inferred scalar implicature rate

## Introduction

The past 15 years have seen the rise and development of a bustling and exciting new field at the intersection of linguistics, psychology, and philosophy: *experimental pragmatics* (Bott & Noveck, 2004; Breheny, Katsos, & Williams, 2006; Degen & Tanenhaus, 2015; Geurts & Pouscoulous, 2009; Grodner, Klein, Carbary, & Tanenhaus, 2010; Huang & Snedeker, 2009; I. A. Noveck & Reboul, 2008) **XXX ADD MORE**. Experimental pragmatics is devoted to experimentally testing theories of how language is used in context. How do listeners draw inferences about the – often underspecified – linguistic signal they receive from speakers? How do speakers choose between the many utterance alternatives they have at their disposal?

The most prominently studied phenomenon in experimental pragmatics is undoubtedly *scalar implicature*. Scalar implicatures arise in virtue of a speaker producing the weaker of two ordered scalemates (Geurts, 2010; Grice, 1975; Hirschberg, 1985; Horn, 1972). Examples are provided in (1-2).

1.

- *Utterance*: Some of her pets are cats.
- *Implicature*: Some, but not all, of her pets are cats.
- *Scale*:

2.

- *Utterance*: She owns a cat or a dog.
- *Implicature*: She owns a cat or a dog, but not both.
- *Scale*:

A listener, upon observing the utterances in (1-2a) typically infers that the speaker intended to convey the meanings in (1-2b), respectively. Since Grice (1975), the agreed-upon

abstract rationalization the listener could give for their inference goes something like this: the speaker could have made a more informative statement by producing the stronger alternative (e.g., *All of her pets are cats.* in (1)). If the stronger alternative is true, they should have produced it to comply with the Cooperative Principle. They chose not to. I believe the speaker knows whether the stronger alternative is true. Hence, it must not be true. The derivation procedure for ad hoc exhaustive inferences such as in (3) is assumed to be calculable in the same way as for scalar implicatures, though the scale is assumed to be contextually driven.

Because the basic reconstruction of the inference is much more easily characterized for scalar implicatures than for other implicatures, scalar implicatures have served as a test bed for many questions in experimental pragmatics, including, but not limited to:

1. Are scalar inferences default inferences, in the sense that they arise unless blocked by (marked) contexts (Degen, 2015; Horn, 1984; Levinson, 2000)?
2. Are scalar inferences default inferences, in the sense that they are computed automatically in online processing and only cancelled by context in a second effortful step if required by context) (Bott & Noveck, 2004; Breheny et al., 2006; Degen & Tanenhaus, 2016; Grodner et al., 2010; Huang & Snedeker, 2009; Politzer-Ahles & Fiorentino, 2013; Tomlinson, Bailey, & Bott, 2013)?
3. What are the (linguistic and extra-linguistic) factors that affect whether a scalar implicature is derived (Bergen & Grodner, 2012; Bonnefon, Feeney, & Villejoubert, 2009; Breheny, Ferguson, & Katsos, 2013; Breheny et al., 2006; Chemla & Spector, 2011; De Neys & Schaeken, 2007; Degen, 2015; Degen & Goodman, 2014; Degen & Tanenhaus, 2015, 2016; Marneffe & Tonhauser, 2016; Potts, Lassiter, Levy, & Frank, 2015; Zondervan, 2010)?
4. How much diversity is there across implicature types, and within scalar implicatures across scale types, in whether or not an implicature is computed (Doran, Ward, Larson,

McNabb, & Baker, 2012; Tiel, Miltenburg, Zevakhina, & Geurts, 2014)?

5. At what age do children acquire the ability to compute implicatures (Barner, Brooks, & Bale, 2011; Horowitz, Schneider, & Frank, 2017; Katsos & Bishop, 2011; Musolino, 2004; Noveck, 2001; Papafragou & Tantalou, 2004; Stiller, Goodman, & Frank, 2015)?

In addressing all of these questions, it has been crucial to obtain estimates of *implicature rates*. For 1., implicature rates from experimental tasks can be taken to inform whether scalar implicatures should be considered default inferences. For 2., processing measures on responses that indicate implicatures can be compared to processing measures on responses that indicate literal interpretations. For 3., contextual effects can be examined by comparing implicature rates across contexts. For 4., implicature rates can be compared across scales (or across implicature types). For 5., implicature rates can be compared across age groups.

A standard measure that has stood proxy for implicature rate across many studies is the proportion of “pragmatic” judgments in truth-value judgment paradigms (Bott & Noveck, 2004; Chemla & Spector, 2011; De Neys & Schaeken, 2007; Degen & Goodman, 2014; Degen & Tanenhaus, 2015; Geurts & Pouscoulous, 2009; Noveck, 2001; Noveck & Posada, 2003). In these kinds of tasks, participants are provided a set of facts, either presented visually or via their own knowledge of the world. They are then asked to judge whether a sentence intended to describe those facts is true or false (or alternatively, whether it is right or wrong, or they are asked whether they agree or disagree with the sentence). The crucial condition for assessing implicature rates in these kinds of studies typically consists of a case where the facts are such that the stronger alternative is true and the target utterance is thus also true but underinformative. For instance, Bott and Noveck (2004) asked participants to judge sentences like “Some elephants are mammals”, when world knowledge dictates that all elephants are mammals. Similarly, Degen and Tanenhaus (2015) asked participants to judge sentences like “You got some of the gumballs” in situations where the visual evidence indicated that the participant received all the gumballs from a gumball

machine. In these kinds of scenarios, the story goes, if a participant responds “FALSE”, that indicates that they computed a scalar implicature, eg to the effect of “Not all elephants are mammals” or “You didn’t get all of the gumballs”, which is (globally or contextually) false. If instead a participant responds “TRUE”, that is taken to indicate that they interpreted the utterance literally as ‘Some, and possibly all, elephants are mammals’ or “You got some, and possibly all, of the gumballs”.

Given the centrality of the theoretical notion of “implicature rate” to much of experimental pragmatics, there is to date a surprising lack of discussion of the basic assumption that it is adequately captured by the proportion of FALSE responses in truth-value judgment tasks (but see Benz and Gotzner (2014); Geurts and Pouscoulous (2009); Degen and Goodman (2014); Katsos and Bishop (2011)). Indeed, the scalar implicature acquisition literature was shaken up when Katsos and Bishop (2011) showed that simply by introducing an additional response option, children started looking much more pragmatic than had been previously observed in a binary judgment paradigm. Katsos and Bishop (2011) allowed children to distribute 1, 2, or 3 strawberries to a puppet depending on “how good the puppet said it”. The result was that children gave on average fewer strawberries to the puppet when he produced underinformative utterances compared to when he produced literally true and pragmatically felicitous utterances, suggesting that children do, in fact, display pragmatic ability even at ages when they had previously appeared not to.

But this raises an important question: in truth-value judgment tasks, how does the researcher know whether an interpretation is literal or the result of an implicature computation? The binary choice task typically used is appealing in part because it allows for a direct mapping from response options – TRUE and FALSE – to interpretations – literal and pragmatic. That the seeming simplicity of this mapping is illusory becomes apparent once a third response option is introduced, as in the Katsos and Bishop (2011) case. How is the researcher to interpret the intermediate option? Katsos and Bishop (2011) grouped the intermediate option with the negative endpoint of the scale for the purpose of categorizing

judgments as literal vs. pragmatic, i.e., they interpreted the intermediate option as pragmatic. But it seems just as plausible that they could have grouped it with the positive endpoint of the scale and taken the hard line that only truly FALSE responses constitute evidence of a full-fledged implicature. The point here is that there has been remarkably little consideration of *linking functions* between behavioral measures and theoretical constructs in experimental pragmatics, a problem in many subfields of psycholinguistics (Tanenhaus, 2004). We argue that it is time to engage more seriously with these issues.

We begin by reporting an experiment that addresses the following question: do the number of response options provided in a truth-value judgment task and the way that responses are grouped into pragmatic (“SI”) and literal (“no SI”) change inferences about scalar implicature rates? Note that this way of asking the question presupposes two things: first, that whatever participants are doing in a truth-value judgment task, the behavioral measure can be interpreted as providing a measure of *interpretation*. And second, that listeners either do or do not compute an implicature on any given occasion. In the General Discussion we will discuss both of these issues. First, following Degen and Goodman (2014), we will offer some remarks on why truth-value judgment tasks are better thought of as measuring participants’ estimates of speakers’ *production* probabilities. This will suggest a completely different class of linking functions. And second, we discuss an alternative conception of scalar implicature as a probabilistic phenomenon, a view that has recently rose to prominence in the subfield of probabilistic pragmatics (Franke & Jäger, 2016; Goodman & Frank, 2016). This alternative conception of scalar implicature, we argue, affords developing and testing quantitative linking functions in a rigorous and motivated way.

Consider a setup in which a listener is presented a card with a depiction of either one or two animals (see Figure 1 for an example). As in a standard truth-value judgment task, the listener then observes an underinformative utterance about this card (e.g., “There is a cat or a dog on the card”) and is asked to provide a judgment on a scale with 2, 3, 4, or 5

response options, with endpoints “wrong” and “right.”<sup>1</sup> In the binary case, this reproduces the standard truth-value judgment task. Figure 1 exemplifies (some of) the researcher’s options for grouping responses. Under what we will call the “Strong link” assumption, only the negative endpoint of the scale is interpreted as evidence for a scalar implicature having been computed. Under the “Weak link” assumption, in contrast, any response that does not correspond to the positive endpoint of the scale is interpreted as evidence for a scalar implicature having been computed. Intermediate grouping schemes are also possible, but these are the ones we will consider here. Note that for the binary case, the Weak and Strong link return the same categorization scheme, but for any number of response options greater than 2, the Weak and Strong link can in principle lead to differences in inferences about implicature rate.




Figure 1. Strong and weak link from response options to researcher inference about scalar implicature rate, exemplified for the disjunctive utterance when the conjunction is true.

Let’s examine an example. Assume three response options (wrong, neither, right). Assume further that a third of participants each gave each of the three responses, i.e., the distributions of responses is 1/3, 1/3, and 1/3. Under the Strong link, we infer that this task yielded an implicature rate of 2/3. Under the Weak link, we infer that this task yielded an implicature rate of 1/3. This is quite a drastic difference if we are for instance interested in

<sup>1</sup>An open question concerns the extent to which the labeling of points on the scale affects judgments (e.g., “wrong”–“right” vs. “false”–“true” vs. “disagree”–“agree”). While some studies have used “false”–“true”, others have argued that judging truth may lead to meta-linguistic reasoning in participants that could distort judgments.



whether scalar implicatures are inference defaults and we would like to interpret an implicature rate of above an arbitrary threshold (e.g., 50%) as evidence for such a claim. Under the Strong link, we would conclude that scalar implicatures are not defaults. Under the Weak link, we would conclude that they are. In the experiment reported in the following section, we presented participants with exactly this setup. We manipulated the number of response options between participants and analyzed the results under different linking assumptions.

## Experiment

Participants played an online card game in which they were asked to judge descriptions of the contents of cards. Different groups of participants were presented with different numbers of response options. On critical trials, participants were presented with descriptions for the cards that typically result in exhaustivity implicatures (“There is a cat on the card” when there was a cat and a dog) or scalar implicatures (“There is a cat or a dog on the card” when there was a cat and a dog). We categorized their responses on such trials according to the weak and the strong link introduced above, and tested whether the number of response options and the linking assumptions led to different conclusions about the rate of computed implicatures in the experimental task.

## Methods

**Participants.** 200 participants were recruited via Amazon Mechanical Turk. They optionally provided demographic information at the end of the study. Participants’ mean age was 35. We also asked participants if they had any prior training in logic. 40 participants reported that they did, while 160 had no prior training in logic. **JD: Do we do any analysis of logic training???** All participants’ data was included in the final analysis.

**Materials and procedure.** The study was administered online through Amazon Mechanical Turk.<sup>2</sup> Participants were first introduced to the set of cards we used in the study

---

<sup>2</sup>The experiment can be viewed [here](#).

(Figure 2). Each card depicted one or two animals, where an animal could be either a cat, a dog, or an elephant. Then participants were introduced to a blindfolded fictional character called Bob. Bob was blindfolded to avoid violations of ignorance expectations associated with the use of disjunction (???; Chierchia et al., 2001) **JD: Masoud, can you insert citation pls? Geurts, Sauerland?**. Participants were told that Bob would guess the contents of the cards and their task was to indicate whether Bob’s guess was wrong or right. On each trial, participants saw a card and a sentence representing Bob’s guess. For example, they saw a card with a cat and read the sentence “There is a cat on the card.” They then provided an assessment of Bob’s guess. The study ended after 24 trials.




Figure 2. Cards used in the connective guessing game.

Two factors were manipulated within participants: card type and guess type. There were two types of cards, cards with only one animal on them and cards with two animals. There were three types of guesses: simple (e.g. *There is a cat*), conjunctive (e.g. *There is a cat and a dog*), and disjunctive (e.g. *There is a cat or a dog*). Crossing card type and guess type yielded trials of varying theoretical interest (see Figure 3): critical underinformative trials that were likely to elicit pragmatic inferences (either scalar or exhaustive) and control trials that were either unambiguously true or false. Each trial type occurred three times with randomly sampled animals and utterances that satisfied the constraint of the trial type. Trial order was randomized.



| elephant                           | cat   | cat and dog                        | cat or dog                                  |   |
|------------------------------------|---|------------------------------------|---|---|
| control:<br>unambiguously<br>false | control:<br>unambiguously<br>true                 | control:<br>unambiguously<br>false | control:<br>unambiguously<br>true           |  |
| control:<br>unambiguously<br>false | <b>critical:<br/>exhaustivity<br/>implicature</b> | control:<br>unambiguously<br>true  | <b>critical:<br/>scalar<br/>implicature</b> |  |

Figure 3. Trial types (critical and control). Headers indicate utterance types. Rows indicate card types. Critical trials are marked in bold.

On critical trials, participants could derive implicatures in two ways. First, on trials on which two animals were present on the card (e.g., cat and dog) but Bob guessed only one of them (e.g. “There is a cat on the card”), the utterance could have a literal interpretation (“There is a cat and possibly another animal on the card”) or an exhaustive interpretation (“There is only a cat on the card”). We refer to these trials as “exhaustive”. Second, on trials on which two animals were on the card (e.g., a cat and a dog) and Bob used a disjunction (e.g., “There is a cat or a dog on the card”), the utterance could have the literal, inclusive, interpretation, or a pragmatic, exclusive interpretation. We refer to these trials as “scalar”.

In order to assess the effect of the number of response options on implicature rate, we manipulated number of response options in the forced choice task between participants. We refer to the choice conditions as “binary” (options: *wrong*, *right*), “ternary” (options: *wrong*, *neither*, *right*), “quaternary” (options: *wrong*, *kinda wrong*, *kinda right*, *right*), and “quinary” (*wrong*, *kinda wrong*, *neither*, *kinda right*, *right*). Thus, the endpoint labels always remained the same. If there was an uneven number of response options, the central option was *neither*. Participants were randomly assigned to one of the four task conditions.

## Results

The collected dataset contains 50 participants in the binary task, 53 in the ternary task, 43 in the quaternary task, and 54 in the quinary task. Figures 4 to 7 show the proportions of response choices in each of the 8 trial types on each of the four response tasks, respectively. We report the relevant patterns of results qualitatively before turning to the quantitative analysis of interest.

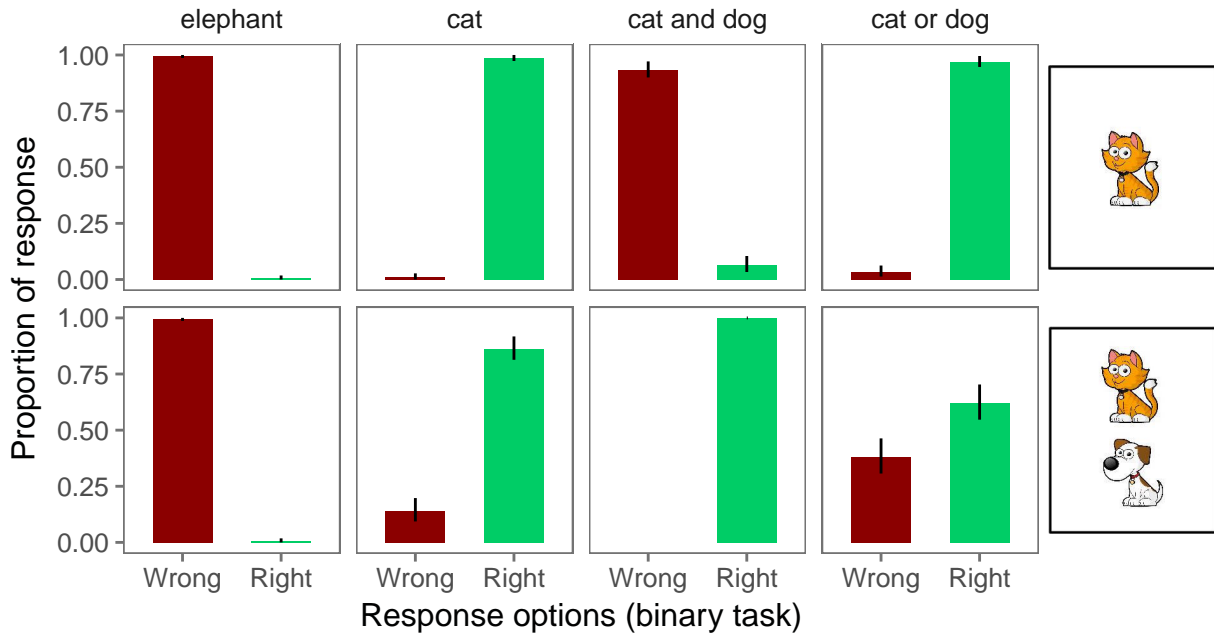


Figure 4. Proportion of responses for the binary forced choice judgments. Error bars indicate 95% bootstrapped confidence intervals.

**Qualitative analysis.** In the binary task, participants were at or close to ceiling in responding “right” and “wrong” on unambiguously true and false trials, respectively (see Figure 4). However, on underinformative trials (“cat” and “cat or dog” description for card with cat and dog), we observe pragmatic behavior: on exhaustive trials, participants judged the utterance “wrong” 14% of the time; on scalar trials, participants judged the utterance “wrong” 38% of the time. That is, both under the weak and the strong link assumptions introduced in the Introduction, inferred implicature rate on exhaustive trials is 14% and on scalar trials 38%.

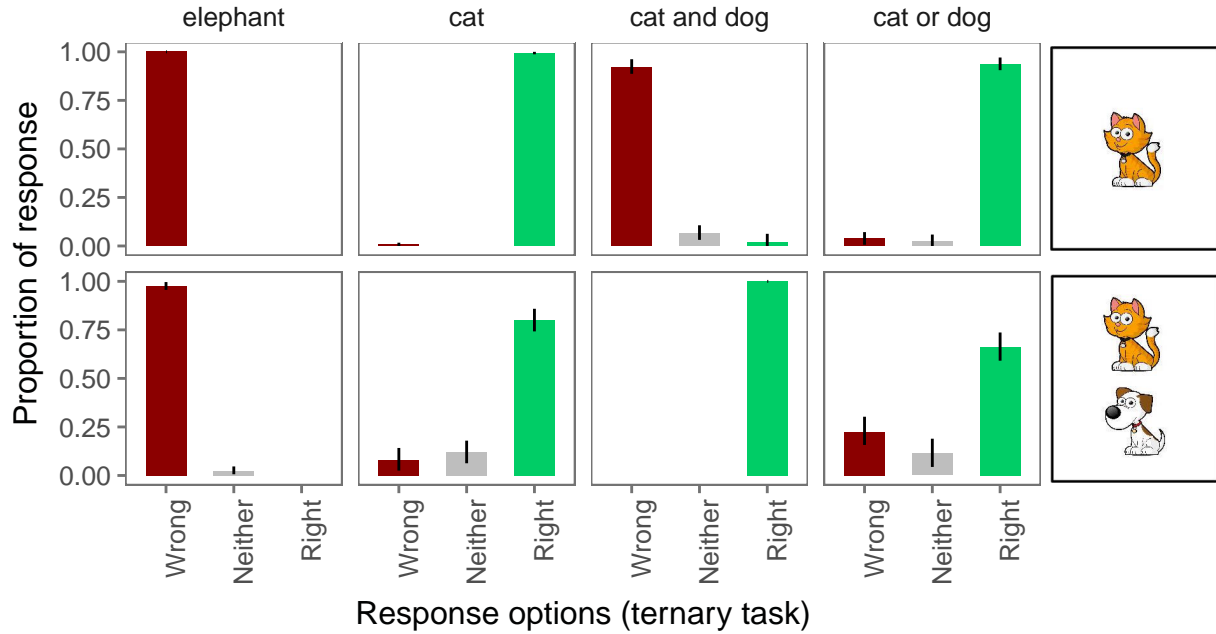


Figure 5. Proportion of responses for the ternary forced choice judgments. Error bars indicate 95% bootstrapped confidence intervals.

In the ternary task, participants were also at or close to ceiling in responding “right” and “wrong” on unambiguously true and false trials, respectively (see Figure 5). And again, on underinformative trials (“cat” and “cat or dog” description for card with cat and dog), we observed pragmatic behavior: on exhaustive trials, participants considered the guess “wrong” 8% of the time and neither wrong nor right 12% of the time. On scalar trials, participants judged the guess “wrong” 23% of the time and “neither” 11% of the time. This means that the weak and strong link lead to different conclusions about implicature rates on the ternary task. Under the weak link, inferred implicature rate on exhaustive trials is 20%; under the strong link it is only 8%. Similarly, under the weak link, inferred implicature rate on scalar trials is 34%; under the strong link it is only 23%.

In the quaternary task (Figure 6), participants were again at or close to ceiling in responding “right” and “wrong” on 4 of the 6 unambiguously true and false trials. However, with four response options, two of the control conditions appear to be showing signs of pragmatic infelicity: when a conjunction was used and only one of the animals was on the

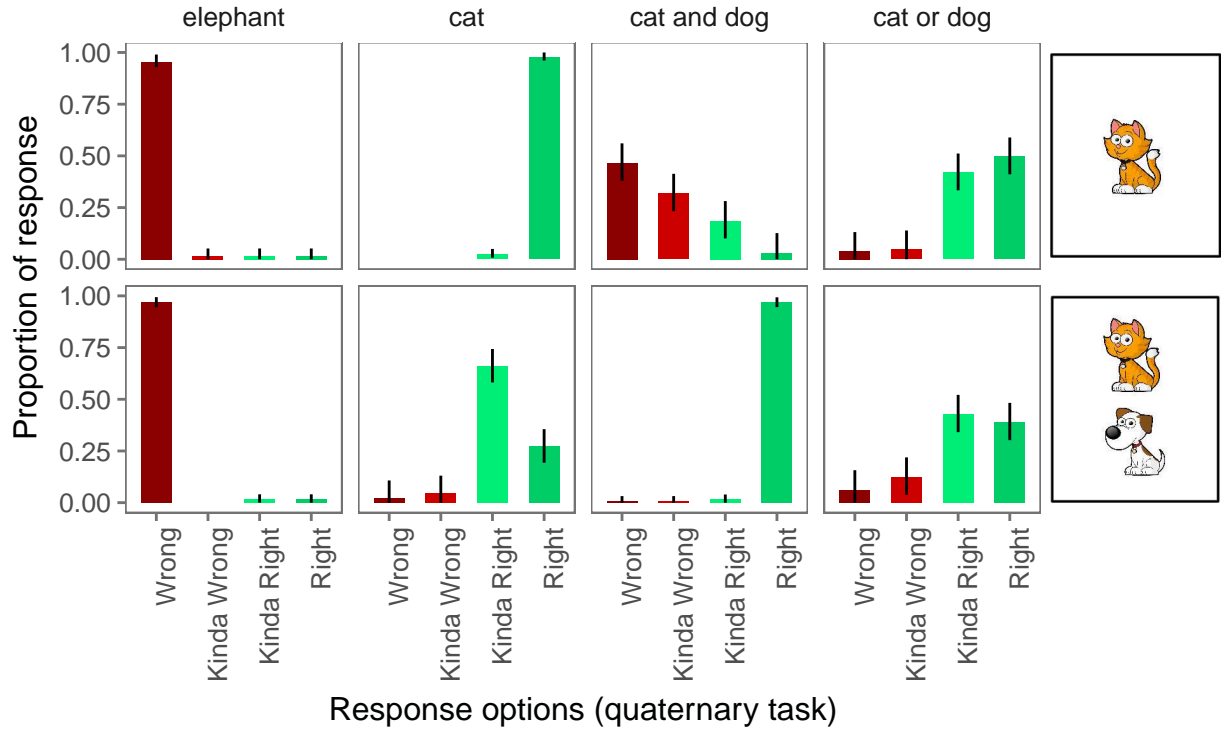


Figure 6. Proportion of responses for the quaternary forced choice judgments. Error bars indicate 95% bootstrapped confidence intervals.

card, participants considered the guess “wrong” most of the time, but they often considered it “kinda wrong” or even “kinda right”. This suggests that perhaps participants considered a notion of partially true or correct statement in our experimental setting. Disjunctive descriptions of cards with only one animal, while previously at ceiling for “right” responses, were downgraded to only “kinda right” 26% of the time, presumably because these utterances are also underinformative, though the degree of underinformativeness may be less egregious than on scalar trials.

On underinformative exhaustive trials, we observed pragmatic behavior as before: participants judged the guess “wrong” 2% of the time, “kinda wrong” 5% of the time, and “kinda right” 66% of the time. On scalar trials, participants judged the guess “wrong” 6% of the time, “kinda wrong” 12% of the time, and “kinda right” 43% of the times.

Thus, we are again forced to draw different conclusions about implicature rates

depending on whether we assume the weak link or the strong link. Under the weak link, inferred implicature rate on exhaustive trials is 73%; under the strong link it is only 2%. Similarly, under the weak link, inferred implicature rate on scalar trials is 61%; under the strong link it is only 6%.

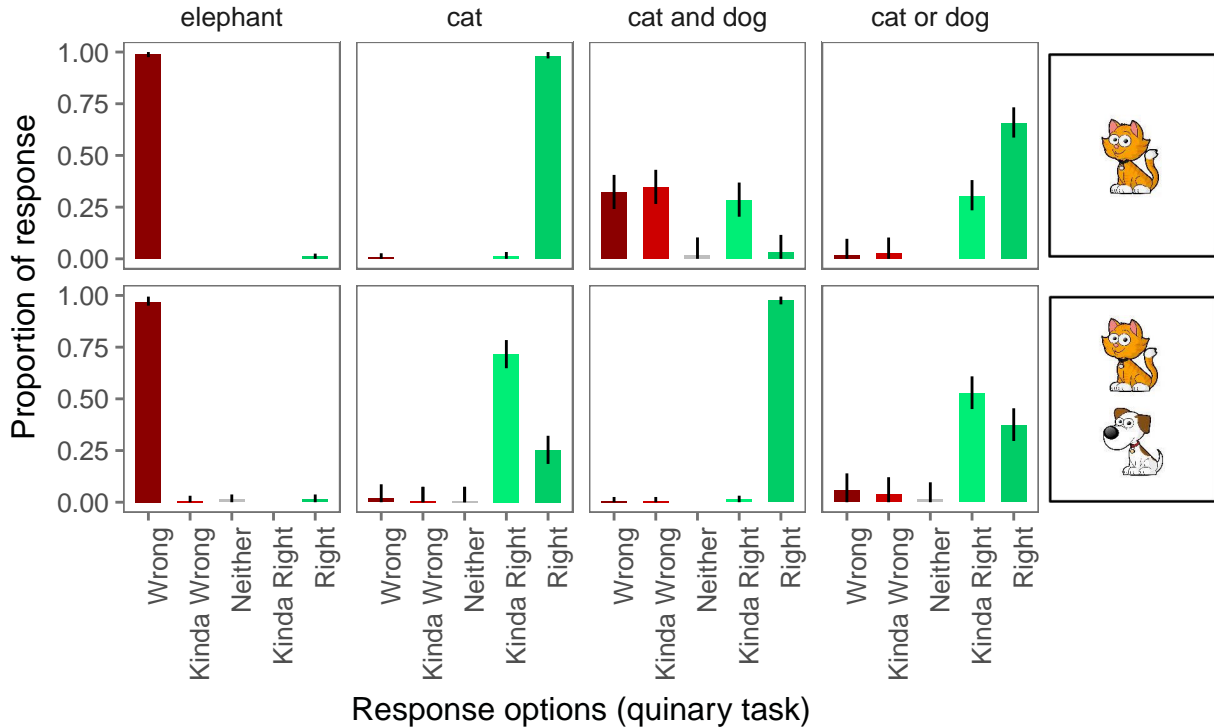


Figure 7. Proportion of responses for the quinary forced choice judgments. Error bars indicate 95% bootstrapped confidence intervals.

Finally, Figure 7 shows the proportion of responses in the quinary task. Performance on the 4 pragmatically felicitous control trials was again at floor and ceiling, respectively. The 2 control conditions that the quaternary task had revealed pragmatic infelicity in again displayed that pragmatic infelicity in the quinary task, suggesting that this is a robust type of pragmatic infelicity that, however, requires fine-grained enough response options to be detected experimentally.

On underinformative exhaustive trials, we observed pragmatic behavior as before: participants judged the guess “wrong” 2% of the time, “kinda wrong” 1 and 1% of the time,

“neither” 1 and 1% of the time, and “kinda right” 72% of the time. On scalar trials, participants judged the guess “wrong” 6% of the time, “kinda wrong” 4% of the time, “neither” 1% of the time, and “kinda right” 52% of the time.

Thus, we would again draw different conclusions about implicature rates depending on whether we assume the weak link or the strong link. Under the weak link, inferred implicature rate on exhaustive trials is 76 and 76%; under the strong link it is only 2%. Similarly, under the weak link, inferred implicature rate on scalar trials is 63%; under the strong link it is only 6%.

**Quantitative analysis.** Our primary goal in this study was to test whether the estimated implicature rate in the experimental task is affected by the linking assumption and the number of response options available to participants. To this end, we only analyzed the critical trials (exhaustive and scalar). In particular, we classified each data point from critical trials as constituting an implicature (1) or not (0) under the strong and weak link. Figure 8 shows the resulting implicature rates by condition and link.

Visually, the weak link tends to result in greater estimates of implicature rates, especially in tasks with more response options. Under the strong link, this latter pattern is reversed: the binary and ternary judgment tasks result in greater estimates of implicature rates than with more response options.

To analyze the effect of link and response options on inferred implicature rate, we used a Bayesian binomial mixed effects model using the R package “brms” (Bürkner & others, 2016) with uninformative or weakly informative priors.<sup>3</sup> The model predicted the log odds of implicature over no implicature from fixed effects of *response type* (binary, ternary, quaternary, quinary – dummy-coded with binary as reference level), *link* (strong vs. weak – dummy-coded with strong as reference level), and trial type (exhaustive vs. scalar – dummy-coded, with exhaustive as reference level), as well as their two-way and three-way interactions. Following Barr, Levy, Scheepers, and Tily (2013), we included the maximal

---

<sup>3</sup>You can access our pre-registration at <https://aspredicted.org/tq3sz.pdf>



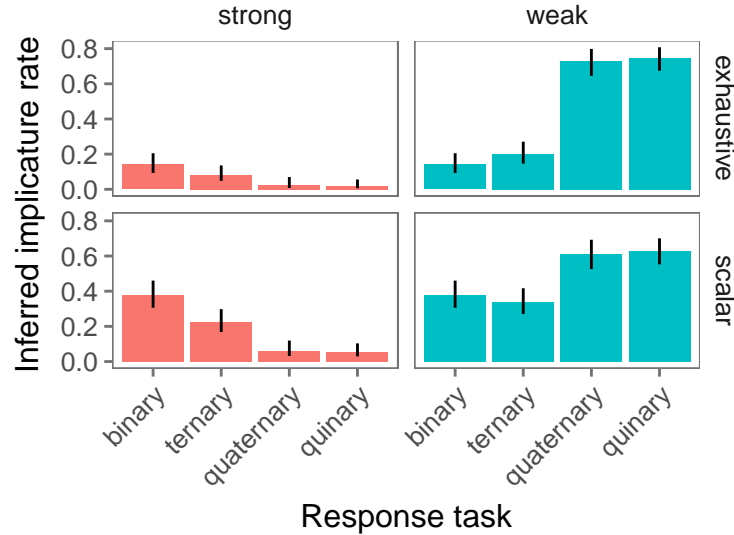


Figure 8. Inferred implicature rates on exhaustive and scalar trials as obtained with the binary, ternary, quaternary, and quinary response task. Columns indicate link from response to implicature rate (strong: proportion of “wrong judgments; weak: proportion of non-’right” judgments).

random effects structure justified by the design: random intercepts for items (cards) and participants, random by-participant slopes for link, trial type, and their interaction, and random by-item slopes for link, trial type, response type, and their interactions. Since the number of response options was a between participant variable we did not include random slopes of response options for participants. Four chains converged after 2000 iterations each (warmup = 1000). Table 1 summarizes the mean parameter estimates and their 95% credible intervals.  $\hat{R} = 1$  for all estimated parameters. All the analytical decisions described here were pre-registered<sup>4</sup>.

[^3: To know more about the default priors of the “brms” package you can refer to [the brms package manual](#).]

The model provided evidence for the following effects: First, there was a main effect of trial type such that scalar trials resulted in greater implicature rates than exhaustive trials

<sup>4</sup>You can access our pre-registration at <https://aspredicted.org/tq3sz.pdf>

(Mean Estimate = 6.09, 95% Credible Interval=[1, 12.29]). Second, there was an interaction between link and number of response options such that the quaternary task (Mean Estimate = 14.03, 95% Credible Interval=[7.24, 21.88]) and the quinary task (Mean Estimate = 17.28, 95% Credible Interval=[10.64, 25.80]) with a weak link resulted in greater implicature rates. Finally, there was a three-way interaction between link, trial type, and number of response options (Mean Estimate = -7.74, 95% Credible Interval=[-16.59, -0.16]). One interpretation of this interaction is that the difference between the weak and strong link on scalar trials in the quinary task was smaller than on exhaustive trials, though we believe this is not too interesting, given that the binary reference level implicature estimate was lower for exhaustive trials in the first place. Crucially, both number of response options and link affect the inferred implicature rate.

## Discussion

We asked whether the linking assumptions and the number of response options available to participants affects the inferred implicature rate in an experimental study. The results presented here suggest they do. A linking assumption that considered the highest point on the scale as literal and any lower point as pragmatic (weak link) resulted in higher implicature rates in tasks with 4 or 5 response options. A linking assumption that considered the lowest point on the scale as pragmatic and any higher point as literal (strong link) reported lower implicature rates in tasks with 4 or 5 options. The results suggest that the choice of linking assumption is a crucial analytical step that can significantly impact the conclusions drawn from truth value judgment tasks. In particular, there is danger for pragmatic ability to be both under- and overestimated.

While the binary truth value judgement task avoids the analytic decision between strong and weak linking assumptions, the results reported here suggest that binary tasks can also underestimate participants' pragmatic competence. In binary tasks, participants are often given the lowest and highest points on a scale ("wrong" vs. "right") and are asked to

report pragmatic infelicities using the lowest point (e.g. “wrong”). The study reported here showed that on trials with true but pragmatically infelicitous descriptions, participants often avoided the lowest point on the scale if they were given more intermediate options. Even though the option “wrong” was available to participants in all tasks, participants in tasks with intermediate options chose it less often. In computing implicature rate, this pattern manifested itself as a decrease in implicature rate under the strong link when more response options were provided, and an increase in implicature rate under the weak link when more response options were provided. These conclusions are in line with Katsos and Bishop (2011)’s argument that pragmatic violations are not as severe as semantic violations and participants do not penalize them as much. Providing participants with only the extreme ends of the scale (e.g. wrong/right, false/true) when pragmatic violations are considered to be of an intermediate nature risks misrepresentation of participants’ pragmatic competence. It further suggests that in studies that use binary tasks to investigate response-contingent processing, proportions of “literal” responses may be a composite of both literal and pragmatic underlying interpretations that just happen to get mapped differently onto different response options by participants.

This study did not investigate the effect of option labels on the inferred implicature rate. However, the results provided suggestive evidence that some options better capture participant intuitions of pragmatic infelicities than others. Among the intermediate options, “kinda right” was chosen most often to report pragmatic infelicities. The option “neither” was rarely used in the ternary and quinary tasks (where it was used as a midpoint), suggesting that participants interpreted pragmatic infelicities as different degrees of being “right” and not “neither right nor wrong.” Therefore, options that capture degrees of being “right” like “kinda right” may prove most suitable for capturing infelicity in the long run. We leave this as a methodological issue for future research.

The study had three further design features worth investigating in future work. First, the utterances were ostensibly produced by a blindfolded character. This was an intentional

decision to control for violation of ignorance expectations with disjunction. A disjunction such as “A or B” often carries an implication or expectation that the speaker is not certain which alternative actually holds. Future work should investigate how the violation of the ignorance expectation interacts with link and number of response options in inferred implicature rate. Second, in this study we considered exhaustive and scalar implicatures with *or*. If the observed effects of link and number of response options hold in general, they should be observable using other scales, e.g., on implicatures with *some*. Finally, our experiment was designed as a guessing game and the exact goal or task-relevant Question Under Discussion of the game was left implicit. Given the past literature on QUD effects on scalar implicature, we expect that different goals – e.g., to help the character win more points vs. to help the character be more accurate – would affect how strict or lenient participants are with their judgments and ultimately affect implicature rate in the task (Degen & Goodman, 2014; Zondervan, 2010). Future work should systematically vary the goal of the game and explore its effects on the inferred implicature rate.

### General Discussion

On the traditional view of the link between implicature and behavior in sentence verification tasks, scalar implicature is conceptualized as a binary, categorical affair - that is, an implicature is either “calculated” or it isn’t, and the behavioral reflexes of this categorical interpretation process should be straightforwardly observed in experimental paradigms. This assumption raises concerns for analyzing variation in behavior on a truth value judgment task; for example, why did the majority of respondents in the binary condition of our experiment answer “Right” to an utterance of the underinformative “There is a cat or dog” when the card had both a cat and a dog on it? And why did a sizeable minority nonetheless choose “Wrong” in this same condition?

To explain these data on the traditional view, we are forced to say that a) not all participants calculated the implicature; or that b) some participants who calculated the

implicature did not choose the anticipated (i.e., “Wrong”) response due to some other cognitive process which overrode the “correct” implicature behavior; or some mixture of (a) and (b). We might similarly posit that one or both of these factors underlie the variation in the ternary, quaternary, and quinary conditions. However, without an understanding of how to quantitatively specify the link between implicature calculation and its behavioral expression, the best we can hope for on this approach is an analysis which predicts general qualitative patterns in the data (e.g. a prediction of relatively more “Right” responses than “Wrong” responses in a given trial of our binary truth-value judgment task, or a prediction of a rise in the rate of response of “Right”/“Wrong” between two experimental conditions, given some contextual manipulation). **XXX Judith, I was thinking of adding the sentence to end this paragraph: “However, we should stress that to the best of our knowledge, even a qualitative analysis of this kind of variation in behavior on sentence verification tasks - much less the effect of the number of response choices on that behavior - is largely underdeveloped scalar implicature literature.” But it might be way too strong a claim?**

**XXX Brandon, can you be more specific about what you mean by “tracing a general qualitative pattern”? XXX Judith: I’ve tried to offer some examples of what I mean by this, see above.**

We contrast the above view of implicature and its behavioral reflexes with an alternative linking hypothesis. Recent developments in the field of probabilistic pragmatics have demonstrated that pragmatic production and comprehension can be captured within the Rational Speech Act (RSA) framework, which treats listeners and speakers as agents that reason about each other and choose interpretations and utterances, respectively, that soft-maximize utility (Bergen, Levy, & Goodman, 2016; Degen, Franke, & Jäger, 2013; Degen, Tessler, & Goodman, 2015; Frank & Goodman, 2012; Franke & Jäger, 2016; Goodman & Frank, 2016; Goodman & Stuhlmüller, 2013; Kao, Wu, Bergen, & Goodman, 2014; Qing & Franke, 2015; Scontras, Degen, & Goodman, 2017). **XXX Judith: re: your**

Slack comments, in the following few sentences I've provided a summary of the intuition behind RSA (going off of the discussion in Goodman & Stuhlmueeller 2013) Much in the spirit of Gricean approaches to pragmatic competence, the RSA framework takes as its point of departure the idea that individuals are rational, goal-oriented communicative agents, who in turn assume that their interlocutors similarly behave according to general principles of cooperativity in communication. Just as in more traditional Gricean pragmatics, pragmatic inference and pragmatically-cooperative language production in the RSA framework are, at their core, the product of counterfactual reasoning about alternative utterances that one might produce (but does not, in the interest of cooperativity). However, the RSA framework explicitly and quantitatively models cooperative interlocutors as agents whose language production and comprehension is a function of Bayesian probabilistic inference regarding other interlocutors' expected behavior in a discourse context.

Specifically, in the RSA framework we model pragmatically competent listeners as continuous probabilistic distributions over possible meanings (states of the world) given an utterance which that listener observes. The probability with which this listener  $L_1$  ascribes a meaning  $s$  to an utterance  $u$  depends upon a prior probability distribution of potential states of the world  $P_w$ , and upon reasoning about the communicative behavior of a speaker  $S_1$ .  $S_1$  in turn is modeled as a continuous probabilistic distribution over possible utterances given an intended state of the world the speaker intends to communicate. This distribution is sensitive to a rationality parameter  $\alpha$ , the production cost  $C$  of potential utterances, and the informativeness of the utterance, quantified via a representation of a literal listener  $L_0$  whose interpretation of an utterance is in turn a function of that utterance's truth conditional content  $[[u]](s)$  and her prior beliefs about the state of the world  $P_w(s)$ .

$$P_{L_1}(s \mid u) \propto P_{S_1}(u \mid s) * P_w(s)$$

$$P_{S_1}(u \mid s) \propto \exp(\alpha(\log(L_0(s \mid u)) - C(u)))$$

$$P_{L_0}(s \mid u) \propto P_w(s) * P_L(s)$$

Thus, our alternative linking hypothesis contrasts with the traditional view in that it is rooted in a quantitative formalization of pragmatic competence which provides us a continuous measure of pragmatic reasoning. In the RSA framework, individuals never categorically draw (or fail to draw) pragmatic inferences about the utterances they hear. For example, exclusivity readings of disjunction are represented in RSA as relatively lower posterior conditional probability of a conjunctive meaning on the  $P_L$  distribution given an utterance of “or”, compared the prior probability of that meaning. Thus, absent auxiliary assumptions about what exactly would constitute “implicature”, it is not even possible to talk about rate of implicature calculation in the RSA framework. The upshot, as we show below, is that this view of pragmatic competence does allow us to talk explicitly and quantitatively about rates of observed behavior in sentence verification tasks.

We take inspiration from the RSA approach and treat participants’ behavior in our experimental tasks as the result of a soft-optimal pragmatic speaker in the RSA framework. That is, following Degen and Goodman (2014), we proceed on the assumption that behavior on sentence verification tasks, such as truth value judgment tasks, is best modeled as a function of an individual’s mental representation of a cooperative interlocutor ( $S_1$  in the language of RSA) rather than of a pragmatic listener who interprets utterances ( $P_{L_1}$ ). In their paper, Degen & Goodman argue that sentence verification tasks are relatively more sensitive to contextual manipulations (such as manipulation of the Question Under Discussion) than are sentence interpretation tasks, and that this follows if sentence interpretation tasks - but not sentence verification tasks - require an additional layer of counterfactual reasoning about the intentions of a cooperative speaker.

A main desideratum of a behavioral linking hypothesis given the RSA view of pragmatic competence is to transform continuous probability distributions into categorical outputs (e.g. responses of “Right”/”Wrong” in the case of the binary condition of our experiment). For a given utterance  $u$  and an intended communicated meaning  $s$ ,  $S_1(u \mid s)$

outputs a conditional probability of  $u$  given  $s$ . For example, in the binary condition of our experiment where a participant evaluated “There is a cat or a dog” when there were both animals on the card, the participant has access to the mental representation of  $S_1$  and hence to the  $S_1$  conditional probability of producing the utterance “cat or dog” given a dog and cat card:  $S_1(\text{“cat or dog”} \mid \text{cat and dog})$ . According to the linking hypothesis advanced here, the participant provides a particular response to  $u$  if the RSA speaker probability of  $u$  lies within a particular probability interval, given an observed state of the world (i.e. the configuration of animals on the card in our experiment). We model a responder,  $R$ , who in the binary condition responds “Right” to an utterance  $u$  in world  $s$  just in case  $S_1(u \mid s)$  exceeds some probability threshold  $\theta$ :

$$\begin{aligned} R(u, w, \theta) \\ &= \text{“Right” iff } S_1(u \mid s) > \theta \\ &= \text{“Wrong” otherwise} \end{aligned}$$

**XXX The following paragraph is mostly new:** Because  $\theta$  is a free parameter in our model, we have a straightforward understanding of between-subject response variation on a given trial of our experiment according to this linking hypothesis. That is, we can attribute this variation to individual differences vis a vis the value of the probability threshold against which utterances are evaluated in a given world state. Moreover, the model of a responder in the binary condition is extended intuitively to the condition where participants had three response options. In this case, we allow for two probability thresholds:  $\theta_1$ , the minimum standard for an utterance in a given world state to count as “Right”, and  $\theta_2$ , the minimum standard for “Neither”. Thus, in the ternary condition,  $R(u, s, \theta_1, \theta_2)$  is “Right” iff  $S_1(u \mid s) > \theta_1$  and “Neither” iff  $\theta_1 > S_1(u \mid s) > \theta_2$ . To fully generalize the model to our five experimental conditions, we say that  $R$  takes as its input an utterance  $u$ , a world state  $s$ , and a number of threshold variables dependent on a variable  $c$ , corresponding to the experimental condition in which the participant finds herself (e.g. the range of possible responses available to  $R$ ).



Given  $c = \text{“ternary”}$

$R(u, w, \theta_1, \theta_2)$

= “Right” iff  $S_1(u | s) > \theta_1$

= “Neither” iff  $\theta_1 > S_1(u | s) > \theta_2$

= “Wrong” otherwise

Given  $c = \text{“quaternary”}$

$R(u, w, \theta_1, \theta_2, \theta_3)$

= “Right” iff  $S_1(u | s) > \theta_1$

= “Kinda Right” iff  $\theta_1 > S_1(u | s) > \theta_2$

= “Kinda Wrong” iff  $\theta_2 > S_1(u | s) > \theta_3$

= “Wrong” otherwise

Given  $c = \text{“quinary”}$

$R(u, w, \theta_1, \theta_2, \theta_3, \theta_4)$

= “Right” iff  $S_1(u | s) > \theta_1$

= “Kinda Right” iff  $\theta_1 > S_1(u | s) > \theta_2$

= “Neither” iff  $\theta_2 > S_1(u | s) > \theta_3$

= “Kinda Wrong” iff  $\theta_3 > S_1(u | s) > \theta_4$

= “Wrong” otherwise

The above analysis is a proof of concept for the following idea: by relaxing the assumptions of the traditional view of scalar implicature (namely, that scalar implicatures either are or are not calculated, and that behavior on sentence verification tasks directly reflects this binary interpretation process), we can propose quantitative models of the variation in behavior we observe in experimental settings.

**XXX The following is entirely new:** However, this way of understanding behavior in truth value judgment tasks also raises new analytic and empirical questions which we must leave to future work. For instance, given the experimental results reported here, what is the range of plausible values for, e.g., the single probability threshold  $\theta$ , which separates

“Right” from “Wrong” utterances in the binary condition? Bayesian statistical methods provide us the means for estimating the value of  $\theta$  in our RSA-based model, by allowing us to recover a posterior distribution of possible values of  $\theta$  given an assumption of a uniform prior and our empirically-observed data, i.e. the actual rate of response of “Right” On each trial of the binary condition of the experiment. We may also ask, for example, whether the speaker probability threshold for an utterance being “Right” is relatively constant across conditions or changes systematically with the introduction of more response options. Again, Bayesian data analysis allows us to explore this question by investigating the extent to which the inferred posterior distribution of, e.g., possible values of the threshold for “Right” in the binary condition of the experiment differs from the inferred posterior distribution of possible values of the threshold for “Right” in the ternary, quaternary, and quinary conditions.

Lastly, the linking analysis proposed here is just one in the space of possible analyses when traditional assumptions about scalar implicature are relaxed. For example, one might reject this threshold-based analysis in favor of one whereby responses are the outcomes of sampling on the (pragmatic speaker or pragmatic listener) probability distributions provided by an RSA model. We again must leave this investigation to future work, but for the time being we emphasize that such quantitative, data-driven and systematic model criticism/comparison is made available to researchers in experimental pragmatics by revising core assumptions about the nature of scalar implicature. Though we no longer have a crisp notion of scalar implicature as something that is or is not “calculated” in interpretation, we have new flexibility to explicitly discuss categorical behavior in experimental settings.

## References

- Barner, D., Brooks, N., & Bale, A. (2011). Accessing the unsaid: the role of scalar alternatives in children’s pragmatic inference. *Cognition*, 118(1), 84–93.  
doi:[10.1016/j.cognition.2010.10.010](https://doi.org/10.1016/j.cognition.2010.10.010)
- Barr, D. J., Levy, R., Scheepers, C., & Tily, H. J. (2013). Random effects structure for

confirmatory hypothesis testing: Keep it maximal. *Journal of Memory and Language*,  
68(3), 255–278.

Benz, A., & Gotzner, N. (2014). Embedded implicatures revisited: Issues with the  
Truth-Value Judgment Paradigm. In J. Degen, M. Franke, & N. D. Goodman (Eds.),  
*Proceedings of the formal & experimental pragmatics workshop, european summer  
school for language, logic and information (esslli)* (pp. 1–6). Tübingen.

Bergen, L., & Grodner, D. J. (2012). Speaker knowledge influences the comprehension of  
pragmatic inferences. *Journal of Experimental Psychology. Learning, Memory, and  
Cognition*, 38(5), 1450–60. doi:[10.1037/a0027850](https://doi.org/10.1037/a0027850)

Bergen, L., Levy, R., & Goodman, N. (2016). Pragmatic reasoning through semantic  
inference. *Semantics and Pragmatics*, 9(1984), 1–46. doi:[10.3765/sp.9.20](https://doi.org/10.3765/sp.9.20)

Bonnefon, J.-F., Feeney, A., & Villejoubert, G. (2009). When some is actually all: scalar  
inferences in face-threatening contexts. *Cognition*, 112(2), 249–58.  
doi:[10.1016/j.cognition.2009.05.005](https://doi.org/10.1016/j.cognition.2009.05.005)

Bott, L., & Noveck, I. (2004). Some utterances are underinformative: The onset and time  
course of scalar inferences. *Journal of Memory and Language*, 51(3), 437–457.  
doi:[10.1016/j.jml.2004.05.006](https://doi.org/10.1016/j.jml.2004.05.006)

Breheny, R., Ferguson, H. J., & Katsos, N. (2013). Taking the epistemic step: Toward a  
model of on-line access to conversational implicatures. *Cognition*, 126(3), 423–40.  
doi:[10.1016/j.cognition.2012.11.012](https://doi.org/10.1016/j.cognition.2012.11.012)

Breheny, R., Katsos, N., & Williams, J. (2006). Are generalised scalar implicatures  
generated by default? An on-line investigation into the role of context in generating  
pragmatic inferences. *Cognition*, 100(3), 434–63. doi:[10.1016/j.cognition.2005.07.003](https://doi.org/10.1016/j.cognition.2005.07.003)

Bürkner, P.-C., & others. (2016). Brms: An r package for bayesian multilevel models using  
stan. *Journal of Statistical Software*, 80(1), 1–28.

Chemla, E., & Spector, B. (2011). Experimental Evidence for Embedded Scalar Implicatures.

576 *Journal of Semantics*, 28(3), 359–400.

577 Chierchia, G., Crain, S., Teresa, M., Guasti, M. T., Gualmini, A., & Meroni, L. (2001). The  
578 Acquisition of Disjunction: Evidence for a Grammatical View of Scalar Implicatures.  
579 In A. H.-J. D. Et al. (Ed.), *BUCLD 25 proceedings* (pp. 157–168). Somerville, MA:  
580 Cascadilla Press.

581 De Neys, W., & Schaeken, W. (2007). When People Are More Logical Under Cognitive Load  
582 - Dual Task Impact on Scalar Implicature. *Experimental Psychology*, 54(2), 128–133.  
583 doi:[10.1027/1618-3169.54.2.128](https://doi.org/10.1027/1618-3169.54.2.128)

584 Degen, J. (2015). Investigating the distribution of 'some' (but not 'all') implicatures using  
585 corpora and web-based methods. *Semantics and Pragmatics*, 8(11), 1–55.  
586 doi:[10.3765/sp.8.11](https://doi.org/10.3765/sp.8.11)

587 Degen, J., & Goodman, N. D. (2014). Lost your marbles? The puzzle of dependent measures  
588 in experimental pragmatics. *Proceedings of the 36th Annual Conference of the*  
589 *Cognitive Science Society*, 397–402.

590 Degen, J., & Tanenhaus, M. K. (2015). Processing scalar implicature A constraint-based  
591 approach. *Cognitive Science*, 39(4), 667–710. doi:[10.1111/cogs.12171](https://doi.org/10.1111/cogs.12171)

592 Degen, J., & Tanenhaus, M. K. (2016). Availability of Alternatives and the Processing of  
593 Scalar Implicatures: A Visual World Eye-Tracking Study. *Cognitive Science*, 40(1),  
594 172–201. doi:[10.1111/cogs.12227](https://doi.org/10.1111/cogs.12227)

595 Degen, J., Franke, M., & Jäger, G. (2013). Cost-Based Pragmatic Inference about  
596 Referential Expressions. *Proceedings of the 35th Annual Conference of the Cognitive*  
597 *Science Society (CogSci'13)*, 376–281.

598 Degen, J., Tessler, M. H., & Goodman, N. D. (2015). Wonky worlds: Listeners revise world  
599 knowledge when utterances are odd. *Proceedings of the 37th Annual Conference of*  
600 *the Cognitive Science Society*, (2), 548–553.

601 Doran, R., Ward, G., Larson, M., McNabb, Y., & Baker, R. E. (2012). A novel experimental  
602 paradigm for distinguishing between what is said and what is implicated. *Language*,

88, 124–154.

Frank, M. C., & Goodman, N. D. (2012). Predicting pragmatic reasoning in language games. *Science*, 336, 998.

Franke, M., & Jäger, G. (2016). Probabilistic pragmatics, or why Bayes' rule is probably important for pragmatics. *Zeitschrift Für Sprachwissenschaft*, 35(1), 3–44.

Geurts, B. (2010). *Quantity implicatures*. Cambridge: Cambridge Univ Press.

Geurts, B., & Pouscoulous, N. (2009). Embedded implicatures?!? *Semantics and Pragmatics*, 2, 1–34. doi:10.3765/sp.2.4

Goodman, N. D., & Frank, M. C. (2016). Pragmatic Language Interpretation as Probabilistic Inference. *Trends in Cognitive Sciences*, 20(11), 818–829. doi:10.1016/j.tics.2016.08.005

Goodman, N. D., & Stuhlmüller, A. (2013). Knowledge and implicature: modeling language understanding as social cognition. *Topics in Cognitive Science*, 5(1), 173–84. doi:10.1111/tops.12007

Grice,

H. P. (1975). Logic and Conversation. *Syntax and Semantics*, 3, 41–58. Retrieved from <http://books.google.com/books?hl=en&lr=&id=hQCzOmaGeVYC&oi=fnd&pg=PA1>

Grodner, D. J., Klein, N. M., Carbary, K. M., & Tanenhaus, M. K. (2010). “Some,” and possibly all, scalar inferences are not delayed: Evidence for immediate pragmatic enrichment. *Cognition*, 116(1), 42–55. doi:10.1016/j.cognition.2010.03.014

Hirschberg, J. (1985). *A Theory of Scalar Implicature* (PhD thesis No. MS-CIS-85-56). University of Pennsylvania; Garland Publishing Company.

Horn, L. (1972). *On the Semantic Properties of the Logical Operators in English* (PhD thesis). UCLA.

Horn, L. (1984). Toward a new taxonomy for pragmatic inference: Q-based and R-based implicature. In D. Schiffrin (Ed.), *Meaning, form, and use in context: Linguistic*

- 629        *applications* (pp. 11–42). Washington: Georgetown University Press.
- 630 Horowitz, A. C., Schneider, R. M., & Frank, M. C. (2017). The Trouble With Quantifiers:  
631        Exploring Children’s Deficits in Scalar Implicature. *Child Development*, 1–40.  
632        doi:[10.1111/cdev.13014](https://doi.org/10.1111/cdev.13014)
- 633 Huang, Y. T., & Snedeker, J. (2009). On-line interpretationf of scalar quantifiers: Insight  
634        into the semantics-pragmatics interface. *Cognitive Psychology*, 58, 376–415.
- 635 Kao, J., Wu, J., Bergen, L., & Goodman, N. D. (2014). Nonliteral understanding of number  
636        words. *Proceedings of the National Academy of Sciences of the United States of*  
637        *America*, 111(33), 12002–12007. doi:[10.1073/pnas.1407479111](https://doi.org/10.1073/pnas.1407479111)
- 638 Katsos, N., & Bishop, D. V. M. (2011). Pragmatic tolerance: implications for the acquisition  
639        of informativeness and implicature. *Cognition*, 120(1), 67–81.  
640        doi:[10.1016/j.cognition.2011.02.015](https://doi.org/10.1016/j.cognition.2011.02.015)
- 641 Levinson, S. C. (2000). *Presumptive Meanings - The Theory of Generalized Conversational*  
642        *Implicature*. MIT Press.
- 643 Marneffe, M.-C. de, & Tonhauser, J. (2016). Inferring meaning from indirect answers to polar  
644        questions: The contribution of the rise-fall-rise contour. In E. Onea, M. Zimmermann,  
645        & K. von Heusinger (Eds.), *Questions in discourse*. Leiden: Brill Publishing.
- 646 Musolino, J. (2004). The semantics and acquisition of number words: integrating linguistic  
647        and developmental perspectives. *Cognition*, 93(1), 1–41.  
648        doi:[10.1016/j.cognition.2003.10.002](https://doi.org/10.1016/j.cognition.2003.10.002)
- 649 Noveck, I. (2001). When children are more logical than adults: experimental investigations  
650        of scalar implicature. *Cognition*, 78(2), 165–188. Retrieved from  
651        <http://www.ncbi.nlm.nih.gov/pubmed/11074249>
- 652 Noveck, I. A., & Reboul, A. (2008). Experimental pragmatics: a Gricean turn in the study  
653        of language. *Trends in Cognitive Sciences*, 12(11), 425–431.  
654        doi:[10.1016/j.tics.2008.07.009](https://doi.org/10.1016/j.tics.2008.07.009)
- 655 Noveck, I., & Posada, A. (2003). Characterizing the Time Course of an Implicature: an

656 Evoked Potentials Study. *Brain and Language*, 85(2), 203–210.

657 doi:[10.1016/S0093-934X\(03\)00053-1](https://doi.org/10.1016/S0093-934X(03)00053-1)

658 Papafragou, A., & Tantalou, N. (2004). Children’s Computation of Implicatures. *Language*  
659 *Acquisition*, 12(1), 71–82.

660 Politzer-Ahles, S., & Fiorentino, R. (2013). The Realization of Scalar Inferences: Context  
661 Sensitivity without Processing Cost. *PLoS ONE*, 8(5).

662 doi:[10.1371/journal.pone.0063943](https://doi.org/10.1371/journal.pone.0063943)

663 Potts, C., Lassiter, D., Levy, R., & Frank, M. C. (2015). Embedded implicatures as  
664 pragmatic inferences under compositional lexical uncertainty. *Journal of Semantics*,  
665 33(1975), 755–802. doi:[10.1093/jos/ffv012](https://doi.org/10.1093/jos/ffv012)

666 Qing, C., & Franke, M. (2015). Variations on a bayesian theme: Comparing bayesian models  
667 of referential reasoning. In H. Zeevat & H.-C. Schmitz (Eds.), *Bayesian natural*  
668 *language semantics and pragmatics* (Vol. 2, pp. 201–220). Cham, Switzerland:  
669 Springer International Publishing. doi:[10.1007/978-3-319-17064-0\\_9](https://doi.org/10.1007/978-3-319-17064-0_9)

670 Scontras, G., Degen, J., & Goodman, N. D. (2017). Subjectivity predicts adjective ordering  
671 preferences. *Open Mind: Discoveries in Cognitive Science*, 1(1), 53–65.

672 doi:[10.1162/opmi](https://doi.org/10.1162/opmi)

673 Stiller, A. J., Goodman, N. D., & Frank, M. C. (2015). Ad-hoc Implicature in Preschool  
674 Children. *Language Learning and Development*, 11(2), 176–190.

675 doi:[10.1080/15475441.2014.927328](https://doi.org/10.1080/15475441.2014.927328)

676 Tanenhaus, M. K. (2004). On-line sentence processing: past, present and future. In M.  
677 Carreiras & C. Clifton (Eds.), *On-line sentence processing: ERPS, eye movements*  
678 *and beyond* (pp. 371–392). London, UK: Psychology Press.

679 Tiel, B. van, Miltenburg, E. van, Zevakhina, N., & Geurts, B. (2014). Scalar diversity.  
680 *Journal of Semantics*. doi:[10.1093/jos/ffu017](https://doi.org/10.1093/jos/ffu017)

681 Tomlinson, J. M., Bailey, T. M., & Bott, L. (2013). Possibly all of that and then some:  
682 Scalar implicatures are understood in two steps. *Journal of Memory and Language*,

683           69(1), 18–35. doi:[10.1016/j.jml.2013.02.003](https://doi.org/10.1016/j.jml.2013.02.003)

684   Zondervan, A. (2010). *Scalar implicatures or focus: an experimental approach* (PhD thesis).

685           Universiteit Utrecht, Amsterdam.



Table 1

*Model parameter estimates and their credible intervals. Rows marked with an asterisk in the evidence column do not contain 0 in the credible interval, thereby providing evidence for an effect.*

| Predictors                                       | Estimate | 2.5%   | 97.5% | Evidence |
|--|----------|--------|-------|----------|
| Intercept  | -8.60    | -13.98 | -4.53 | *        |
| Link = Weak                                      | -0.15    | -4.86  | 4.77  |          |
| Task = Quaternary                                | -1.83    | -8.08  | 4.20  |          |
| Task = Quinary                                   | -4.05    | -10.90 | 2.38  |          |
| Task = Ternary                                   | -1.45    | -7.31  | 4.56  |          |
| Implicature = Scalar                             | 6.09     | 1.00   | 12.29 | *        |
| Link = Weak : Task = Quaternary                  | 14.03    | 7.24   | 21.88 | *        |
| Link = Weak : Task = Quinary                     | 17.28    | 10.64  | 25.80 | *        |
| Link = Weak : Task = Ternary                     | 3.81     | -1.49  | 9.22  |          |
| Link = Weak : Implicature = Scalar               | 0.90     | -4.01  | 6.43  |          |
| Task = Quaternary : Implicature = Scalar         | -5.67    | -13.66 | 1.54  |          |
| Task = Quinary : Implicature = Scalar            | -2.31    | -9.30  | 4.61  |          |
| Task = Ternary : Implicature = Scalar            | -1.31    | -7.70  | 4.65  |          |
| Link=Weak : Task=Quaternary : Implicature=Scalar | -3.29    | -12.07 | 4.55  |          |
| Link=Weak : Task=Quinary : Implicature=Scalar    | -7.74    | -16.59 | -0.16 | *        |
| Link=Weak : Task=Ternary : Implicature=Scalar    | -1.44    | -7.00  | 4.22  |          |